

Pedro Henrique Rocha Moy

AI, Machine Learning, Data Science & Engineering Miami, FL

Education

Executive MBA in Business Administration

University of Miami - Miami, 2021 - 2022

Master's Degree in Computer Science (Machine Learning)

Georgia Institute of Technology - Atlanta, GA, 2015 - 2017

Master's Degree in Earth Science and Engineering (Geophysics)

King Abdullah University of Science and Technology - Saudi Arabia, 2010 - 2012

Bachelor's Degree in Mechanical Engineering (Math Minor)

University of Massachusetts Lowell - Lowell, MA, 2008 - 2010

Experience

Chief Architect | Rocha Moy Trading & Capital

2017 - Present

- Developed the API for probabilistic and algorithmic options trading with Interactive Brokers and TD Ameritrade
- Specialties include data integration, task automation, portfolio simulations, risk mitigation, and strategy validation
- Integrated many different data sources from APIs to web scraping
- Automated trade execution, trade scheduling, and fund releases for trading
- Technologies: Python, AWS, REST APIs, Web Scraping, Machine Learning, Probability Theory, Statistical Modeling, Time Series Analysis, Quantitative Finance, Options Trading, Portfolio Optimization, Risk Management, Backtesting, Simulations, Data Integration, Data Pipelines, Task Automation

AI Consultant | Toptal Client

2025

- Developed a prototype platform that automated LLC formation workflows, including a detailed regulatory study for Florida, reducing manual research time for new founders by well over 50%

- Built a back-end LLM system that validated user inputs and flagged compliance gaps during the LLC creation process, ensuring accuracy of filings
- Designed a domain-specific knowledge base and implemented an RAG pipeline to deliver precise, context-aware guidance to users
- Technologies: Python, LLMs, RAG, NLP, Prompt Engineering, Vector Databases, Knowledge Base Design, Document Processing, Regulatory Compliance, Prototype Development

AI Software Developer | Toptal Client

2023 - 2024

- Built a full-stack Django and React application from scratch that parsed contract requirement PDFs into structured suggestions and clarification questions, enabling teams to accelerate project scoping significantly
- Engineered an autoscaling cloud infrastructure using ECS and IaC best practices to support stable, high-throughput processing of large procurement documents, ensuring near-zero downtime
- Designed the end-to-end LLM architecture with OpenAI models and served as the primary expert for all LLM-related issues, improving extraction accuracy and capabilities, shaping model-driven workflows across both the Python and JavaScript stacks
- Technologies: Python, JavaScript, Django, React, OpenAI API, LLMs, NLP, Prompt Engineering, AWS ECS, Infrastructure as Code, Docker, PDF Processing, Autoscaling, Full-stack Development

Lead Data Scientist | Deloitte Contractor

2021 - 2022

- Designed, implemented, and deployed different natural language processing models
- Worked with stakeholders to understand use cases, the pathway to product development, and implementation using deployed models
- Mentored and supported junior data scientists on the team
- Technologies: Python, NLP, Machine Learning, Model Deployment, Stakeholder Management, Data Analysis, Mentoring

Enterprise Lead Data Architect | Toptal Client

2020 - 2022

- Handled the architecture, development, and automation of distributed computing pipelines and data storage in the cloud for the enterprise
- Automated scalable infrastructure in the cloud to respond to development and consumer demand

- Co-managed and supervised a team of engineers from designing and delegating tasks, mentoring, and overseeing work
- Technologies: Python, AWS, Amazon EMR, Spark, Distributed Computing, Data Lake Architecture, Cloud Infrastructure, Data Pipelines, Data Storage, Team Leadership

Enterprise Senior ETL and Data Engineer | Toptal Client

2019 - 2020

- Designed, implemented, and deployed to production fully-fledged distributed ETL jobs in Spark/Scala API
- Worked with various sources and sinks of data including disparate files, Hive tables, Mongo collections, and Kafka brokers
- Served as the senior engineer and tech lead of the team strengthening engineering and development processes, improving software quality control, and helping design stories for sprints
- Technologies: Scala, Python, Spark, MongoDB, Hive, Kafka, Distributed Computing, ETL, Data Pipelines, Agile/Scrum, Code Review, Tech Leadership

Hadoop Proof of Concept for Atmospheric Sciences Project | Toptal Client

2019 - 2020

- Built a cluster from scratch, adhering to the client's needs to work with the home cluster
- Designed and implemented generic and specific data architectures meeting the client's query complexity and performance needs
- Built PySpark and Python software layers of abstraction to allow the client to build on top of the current infrastructure
- Technologies: PySpark, Python, Hadoop, HDFS, Big Data, Data Architecture, Cluster Management, Atmospheric Sciences

Research Data Engineer | Nicklaus Children's Hospital

2018 - 2019

- Developed existing analytical and data workflows for users of R, Python, and Impala establishing best engineering practices
- Provided ad hoc and systematically developed ETL and big data pipelines, validation, and integration of varying data sources
- Liaised for the research department to IT and BI departments providing guidance and expertise on analytical and data needs
- Technologies: Spark, Scala, Python, R, Impala, ETL, Big Data, Healthcare Data, Data Validation, Business Intelligence

Technical Advisor | Insight Data Science

2018

- Worked with fellows and their data engineering projects on problem definition, systems architecture, and execution
- Advised on technologies such as Spark, Kafka, Redis, HBase, Cassandra, and PostgreSQL
- Conducted mock interviews with fellows on scalability concepts, algorithms, and CS fundamentals
- Technologies: Spark, Kafka, Redis, HBase, Cassandra, PostgreSQL, Distributed Systems, Systems Architecture, Mentoring, Algorithms

Senior Software Engineer | NexHealth

2016 - 2017

- Developed and deployed software to the client's site to perform data collection and server sync
- Performed both database and web-based data integrations of electronic medical records back to NexHealth servers
- Developed a smart SMS response system allowing the user to interact with NexHealth products via SMS
- Technologies: Python, Scala, JavaScript, Redis, PostgreSQL, Apache Spark, REST APIs, Data Integration, SMS/Messaging, Electronic Health Records, Web Development

Data Scientist | QuaEra Insights

2016

- Served as the lead data scientist in a consulting project overseeing data management and modeling strategy
- Used natural language processing to transform unstructured data into features and extract business intelligence
- Built a recommendation engine as business rules potentially yielding savings on up to 50% of the business
- Technologies: Python, NLP, Machine Learning, Recommendation Systems, Feature Engineering, Business Intelligence, Data Management

Data Engineering Fellow | Insight Data Science

2015

- Built themidgame-tube, a platform designed to discover YouTube influencers on brand names worldwide
- Deployed Amazon's EMR Spark with HBase processing and ingesting billions of data tuples
- Attained linear scalability performance tested with up to 20 nodes

- Technologies: Python, AWS, Apache Spark, EMR, HBase, Big Data, Distributed Computing, Data Ingestion, Scalability Testing

Data Analyst | Cartesian

2015

- Aided managed analytics efforts promoting best practices within batch workflows and data management
- Conducted independent research into big data workflows considering data mining and BI integration
- Built short data pipelines consuming APIs, transforming, loading, and exposing data connections to BI tools
- Technologies: Python, PostgreSQL, REST APIs, ETL, Data Pipelines, Business Intelligence, Data Mining, Data Analytics

Data Analytics Engineer | Daktari Diagnostics

2013 - 2015

- Worked as the lead developer of mainstream data processing and data analysis applications in Python for Windows/Mac
- Developed a calibration model for the Daktari CD4 testing device improving the system's accuracy by 20-30%
- Deployed machine learning models embedded in standalone applications to end users for data classification
- Technologies: Python, Machine Learning, Statistical Modeling, Calibration, Data Classification, Desktop Applications, Medical Devices, Data Processing

Portfolio

Pastoral Conscience AI

https://rocha-moy-engineering-technology.github.io/pastoral_conscience_site/

Built an “Artificial Conscience” AI system that generates scripture-grounded spiritual reflections using RAG retrieval from Psalms and DSPy-based conscience checkers. Designed and implemented a governed reasoning pipeline with three verification layers: Helpfulness, Psalm Grounding, and Citation Integrity checks. The system enforces strict alignment of beliefs; every response is scored, verified, and corrected before delivery. Architecture: Go back end with hexagonal architecture, Gemini File Search API for RAG retrieval, Python and DSPy services for conscience checkers with model fallback, Next.js front end with real-time SSE progress updates, and PostgreSQL for data persistence.